

# **APPLIED DATA SCIENCE GROUP-2**

## **IMDB SCORE PREDICTION (PHASE - 5)**

### **DEVELOPMENT & SUBMISSION**

#### **PROBLEM STATEMENT**

The problem statement for the IMDb score prediction project:

"In the film industry, the success of a movie is often measured by its rating. The Internet Movie Database (IMDb) is one of the most popular platforms where movies are rated by viewers. These ratings play a crucial role in influencing audience viewership and are of great interest to film production companies, directors, and investors.

The task at hand is to build a predictive model that can accurately estimate the IMDb score of a movie before its release based on various factors such as director, actors, budget, genres, etc. This would provide valuable insights into the potential success of a movie and could guide decision-making processes in film production and marketing strategies.

The challenge lies in selecting relevant features from the available data, handling missing or inconsistent data, choosing an appropriate regression algorithm for prediction, and validating the model's performance using suitable metrics. The goal is to achieve a model with high accuracy and robustness that can generalize well to new, unseen movie data."

#### **DESIGN AND ANALYSIS**

The design and analysis for the IMDb score prediction project can be broken down into several steps:

##### **1. Problem Understanding:**

The first step is to understand the problem at hand. The goal is to predict the IMDb score of a movie based on various features. This is a regression problem as the IMDb score is a continuous variable.

## **2. Data Collection:**

The next step is to collect data that can be used to train the model. The dataset used in this project contains movie metadata from IMDb, including features such as director, actors, budget, genres, etc.

## **3. Exploratory Data Analysis (EDA):**

This involves understanding the dataset through descriptive statistics and visualizations. It helps in identifying patterns, correlations, and outliers in the data.

## **4. Data Preprocessing:**

This step involves cleaning the data and making it suitable for modeling. It includes handling missing values, encoding categorical variables, feature scaling, etc.

## **5. Feature Selection:**

Not all features in the dataset may be useful for predicting the IMDb score. Feature selection methods can be used to select the most relevant features.

## **6. Model Building:**

This involves choosing a suitable regression algorithm and training it on the preprocessed data. The choice of algorithm depends on the nature of the data and the problem.

## **7. Model Evaluation:**

After training the model, it's important to evaluate its performance using suitable metrics. For regression problems, metrics like Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and R-squared can be used.

## **8. Model Optimization:**

Based on the evaluation results, the model may need to be optimized by tuning its hyperparameters or using different algorithms.

## **9. Model Deployment:**

Once the model is optimized and achieves satisfactory performance, it can be deployed for use in predicting IMDb scores of new movies. Throughout these steps, it's important to document each step and decision made for future reference and reproducibility.

### **Phase-1:**

#### **Problem Definition and Design Thinking**

- Introduction
- Primary goals
- Sample dataset
- Problem
- Design & Analysis
- Summary

### **Phase-2: Innovation**

Consider exploring advanced regression techniques like Gradient Boosting or Neural Networks for improved prediction accuracy

- Introduction
- Dataset details
- Libraries
- How to train and test
- Process flow

### **Phase 3 – Development part -1**

Begin building the IMDb score prediction model by loading and preprocessing the dataset.

- Loading the dataset
- Preprocessing the dataset

## Phase 4 – Development part -2

Continue building the IMDb score prediction model by feature engineering, model training, and evaluation

- Model training
- Feature engineering
- Evaluation

## Phase 5 – Project Documentation & Submission

Document the IMDb score prediction project and prepare it for submission.

- Documentation
- Submission

## DATASET

dataset: <http://www.kaggle.com>

dataset name: IMDB score Prediction

dataset link: <https://www.kaggle.com/datasets/bobirino/movie-metadata>

Detail	Compact	Column	10 of 28 columns			
▲ color		▲ director_name	# num_critc_for_re...	# duration	# director_faceboo...	4
Color	95%	[null]	2%			
Black and White	4%	Steven Spielberg	1%			
Other (19)	0%	Other (4913)	97%	1 813	7 511	0 23.0k C
Color		Gore Verbinski	302	169	563	1
Color		Sam Mendes	602	148	0	1
Color		Christopher Nolan	813	164	22000	2
		Doug Walker			131	
Color		Andrew Stanton	462	132	475	5
Color		Sam Raimi	392	156	0	4
Color		Nathan Greno	324	100	15	2
Color		Joss Whedon	635	141	0	1

## STEPS FOR PREPROCESSING DATA:

1. **Dataset Acquisition:** The initial step is to gather the dataset that you'll be working with.
2. **Library Importation:** Import all necessary libraries that will be used for data preprocessing.
3. **Dataset Importation:** Load the dataset into your working environment.
4. **Handling Missing Values:** Identify any missing values in the dataset and decide on the best strategy to handle them, such as imputation or deletion.
5. **Categorical Data Encoding:** Convert categorical data into a format that can be easily understood by machine learning algorithms, typically through one-hot encoding or label encoding.
6. **Dataset Division:** Split the dataset into a training set and a testing set. This allows for proper evaluation of the model's performance.
7. **Feature Scaling:** Normalize or standardize features to ensure that they're on a similar scale and to prevent any one feature from dominating others.

### 1.Acquire the dataset:

**Dataset link:** <https://www.kaggle.com/datasets/bobirino/movie-metadata>

color	director_n	num_critic	duration	director_fr	actor_3	actor_2_n	actor_1_f	gross	genres	actor_1_n	movie_title	num_votes	cast_total	actor_3_n	face	num	plot_keyw	movie_im	num_user	language	country	content_r	budget
Color	James Can	723	178	0	855	Joel David	1000	760505847	Action Ad CCH Poun	Avatar	886204	4834	Wes Studi	0	avatar fut	http://www	3054	English	USA	PG-13	237000		
Color	Gore Verbi	302	169	563	1000	Orlando Bl	40000	309404152	Action Ad Johnny De	Pirates of	471220	48350	Jack Daver	0	goddess n	http://www	1238	English	USA	PG-13	300000		
Color	Sam Mend	602	148	0	161	Rory Kinne	11000	200074175	Action Ad Christoph	Spectre	275868	11700	Stephanie	1	bomb esp	http://www	994	English	UK	PG-13	245000		
Color	Christophe	813	164	22000	23000	Christian B	27000	448130642	Action Th Tom Hard	The Dark k	1144337	106759	Joseph Go	0	deception	http://www	2701	English	USA	PG-13	250000		
	Doug Walker			131		Rob Walke	131		Document Doug Wall	Star Wars:	8	143		0		http://www.imdb.com/title/tt5289954/?ref_=fn_tt_tt_1							
Color	Andrew St	462	132	475	530	Samantha	640	73058679	Action Ad Daryl Saba	John Carte	212204	1873	Polly Walk	1	alien ame	http://www	738	English	USA	PG-13	263700		
Color	Sam Raimi	392	156	0	4000	James Frar	24000	336530303	Action Ad J.K. Simmc	Spider-Ma	383056	46055	Kirsten Dui	0	sandman	http://www	1902	English	USA	PG-13	258000		
Color	Nathan Gr	324	100	15	284	Donna Mu	799	200807262	Adventure Brad Garre	Tangled	294810	2036	M.C. Gain	1	17th centu	http://www	387	English	USA	PG	260000		
Color	Joss Whed	635	141	0	19000	Robert Do	26000	458991599	Action Ad Chris Hem	Avengers:	462669	92000	Scarlett Jo	4	artificial	in http://www	1117	English	USA	PG-13	250000		
Color	David Yate	375	153	282	10000	Daniel Rad	25000	301956980	Adventure Alan Rickn	Harry Pott	321795	58753	Rupert Gri	3	blood boc	http://www	973	English	UK	PG	250000		
Color	Zack Snyder	673	183	0	2000	Lauren Col	15000	330249062	Action Ad Henry Cavi	Batman v	371639	24450	Alan D. Pui	0	based on c	http://www	3018	English	USA	PG-13	250000		
Color	Bryan Sing	434	169	0	903	Marlon Bri	18000	200069408	Action Ad Kevin Spac	Superman	240396	29991	Frank Lang	0	crystal epi	http://www	2367	English	USA	PG-13	209000		
Color	Marc Forst	403	106	395	393	Mathieu A	451	168368427	Action Ad Giancarlo	Quantum c	330784	2023	Rory Kinne	1	action her	http://www	1243	English	UK	PG-13	200000		
Color	Gore Verbi	313	151	563	1000	Orlando Bl	40000	423032628	Action Ad Johnny De	Pirates of	522040	48486	Jack Daver	2	box office	http://www	1832	English	USA	PG-13	225000		
Color	Gore Verbi	450	150	563	1000	Ruth Wils	40000	89289910	Action Ad Johnny De	The Lone f	181792	45757	Tom Wilki	1	horse out	http://www	711	English	USA	PG-13	215000		
Color	Zack Snyder	733	143	0	748	Christophe	15000	291021565	Action Ad Henry Cavi	Man of Ste	548573	20495	Harry Lenr	0	based on c	http://www	2536	English	USA	PG-13	225000		
Color	Andrew Ac	258	150	80	201	Pierfrance	22000	141614023	Action Ad Peter Dink	The Chron	149922	22697	Dami	4	brother br	http://www	438	English	USA	PG	225000		
Color	Joss Whed	703	173	0	19000	Robert Do	26000	623279547	Action Ad Chris Hem	The Aveng	995415	87697	Scarlett Jo	3	alien invas	http://www	1722	English	USA	PG-13	220000		
Color	Rob Mars	448	136	252	1000	Sam Claflir	40000	241063875	Action Ad Johnny De	Pirates of	370704	54083	Stephen Gi	4	blackbear	http://www	484	English	USA	PG-13	250000		
Color	Barry Sonr	451	106	188	718	Michael St	10000	179020854	Action Ad Will Smith	Men in Bla	268154	12572	Nicole Sch	1	alien crim	http://www	341	English	USA	PG-13	225000		
Color	Peter Jack	422	164	0	773	Adam Brov	5000	255108370	Adventure Aidan Turri	The Hobbi	354228	9152	James Nes	0	army elf	http://www	802	English	New Zeala	PG-13	250000		
Color	Marc Web	599	153	464	963	Andrew Ge	15000	262030663	Action Ad Emma Sto	The Amazi	451803	28489	Chris Zylka	0	lizard out	http://www	1225	English	USA	PG-13	230000		
Color	Ridley Sco	343	156	0	738	William Hl	891	105219735	Action Ad Mark Addy	Robin Hoo	211765	3244	Scott Grim	0	1190s arc	http://www	546	English	USA	PG-13	200000		
Color	Peter Jack	509	186	0	773	Adam Brov	5000	258355354	Adventure Aidan Turri	The Hobbi	483540	9152	James Nes	6	dwarf elf	http://www	951	English	USA	PG-13	225000		
Color	Chris Weit	251	113	129	1000	Eva Green	16000	70083519	Adventure Christophe	The Golde	149019	24106	Kristin Sco	2	children e	http://www	666	English	USA	PG-13	180000		

## DESCRIBE DATA USED

To predict IMDb scores, a wide variety of data can be used. The accuracy and effectiveness of the prediction can depend on the complexity and size of the dataset, as well as the specific machine learning or statistical techniques used.

**1. Movie Metadata:** This includes information like the movie's title, release date, genre, director, writer, and production company.

**2.Cast and Crew Information:** Data about the actors, actresses, and key crew members, such as the director, producer, and writer.

**3. Box Office Performance:** Details about a movie's financial success, including its budget, gross revenue, and profitability.

**4. Movie Plots and Summaries:** Descriptions of the movie's plot or summary, which can be used for natural language processing and sentiment analysis.

**5. User Reviews and Ratings:** IMDb has user-generated reviews and ratings, which can be used as features. Sentiment analysis on user reviews can provide valuable insights.

**6. Movie Posters and Trailers:** Visual data like movie posters and trailers can be analyzed for aesthetic and marketing factors.

**7. Awards and Nominations:** Information about awards won or nominated for can indicate a movie's quality and critical reception.

**8.Runtime:** The duration of the movie can sometimes be a predictor of IMDb score, as very long or very short movies may have different audience expectations.

**9.Release Information:** Data about the geographical and temporal release of the movie can be relevant, as release strategies vary.

**10.Social Media Data:** Information about a movie's presence on social media platforms, such as the number of followers on official pages, can be used to gauge audience engagement.

## **Innovation & Design:**

### **• Personalized Recommendations:**

Develop a recommendation system that suggests movies based on a user's IMDb ratings and preferences. Use machine learning algorithms to make these recommendations more accurate over time.

### **• Visual Data Analytics:**

Create interactive visualizations that allow users to explore IMDb data. Visualize trends, ratings distribution, and other insights that go beyond a simple numeric score.

### **• Crowdsourced Reviews:**

Enable users to submit detailed reviews alongside their ratings. Implement sentiment analysis to summarize the overall sentiment of reviews.

### **• Filter and Search Enhancement:**

Improve IMDb's filtering and search capabilities. Add advanced filters such as genre-specific, release year, or director-based searches.

### **• IMDb Score Predictions:**

Develop a model to predict IMDb scores for movies before they are released, considering factors like the cast, crew, and pre-release buzz.

### **• User-Generated Lists:**

Allow users to create and share lists of their favorite movies, creating a sense of community and enabling users to discover new films.

### **• Mobile App Integration:**

Create a mobile app that seamlessly integrates these features, making it easy for users to access IMDb's enhanced functionalities on the go.

### **• Data Insights for Filmmakers:**

Offer insights to filmmakers and studios on how their movies are rated by users, potentially helping them make improvements in future projects.

- **Rating Aggregation:**

Aggregate IMDb ratings with ratings from other sources like Rotten Tomatoes or Metacritic to provide a more comprehensive view of a movie's reception.

- **Accessibility and User-Centric Design:**

Ensure the platform is accessible to all users, including those with disabilities, and prioritize a user-centric design for a seamless experience.

- **Community Engagement:**

Implement features like forums or discussion boards where users can engage in meaningful discussions about movies and ratings.

- **Data Security and Privacy:**

Pay strict attention to data security and user privacy, especially when handling user-generated content and personal preferences.

## **Details of Libraries:**

- **IMDbPY (Python Library):**

IMDbPY is a Python package specifically designed to access and retrieve data from the IMDb website. You can use it to fetch movie details, cast information, user reviews, and IMDb ratings.

- **OMDb API:**

The Open Movie Database (OMDb) API provides access to a vast amount of movie-related data, including IMDb ratings, plot summaries, release dates, and more. It's easy to use and doesn't require an IMDb API key.

- **TMDb API:**

The Movie Database (TMDb) API offers movie information, including user ratings and reviews. While not IMDb-specific, it can complement IMDb data for a broader perspective.



- **Beautiful Soup (Python Library):**

Beautiful Soup is a Python library for web scraping. You can use it to extract data from IMDb web pages when IMDbPY doesn't provide the specific data you need.

- **Pandas (Python Library):**

Pandas is a powerful data manipulation and analysis library for Python. It's great for handling, cleaning, and analyzing IMDb data obtained through IMDbPY or web scraping.

- **Matplotlib and Seaborn (Python Libraries):**

Matplotlib and Seaborn are popular Python libraries for data visualization. You can use them to create various charts and plots to visualize IMDb ratings distribution, trends, and correlations.

- **Scikit-learn (Python Library):**

If your project involves machine learning or predictive modeling based on IMDb data, Scikit-learn is a go-to library for tasks like building recommendation systems or predicting IMDb scores.

- **D3.js (JavaScript Library):**

If you plan to create interactive web visualizations for IMDb data, D3.js is a powerful JavaScript library for data-driven graphics. It can be used to create dynamic and interactive data visualizations.

- **SQLite (Database):**

SQLite is a lightweight database engine that you can use to store and manage IMDb-related data locally, making it easier to query and analyze large datasets.

- **Flask (Python Framework):**

If you're building a web application to present IMDb-related data, Flask is a lightweight Python web framework that can help you develop the backend of your application.

- **React or Vue.js (JavaScript Frameworks):**

For the frontend of a web application, React or Vue.js can be useful frameworks to create responsive and interactive user interfaces.

- **Heroku or AWS (Cloud Services):**

If you plan to deploy your IMDb-related project online, platforms like Heroku or AWS can host your web application and databases in the cloud.

## **EVALUATION AND METRICS**

When building a model for IMDb score prediction, it's important to evaluate its performance to determine how well it predicts IMDb scores. Here are common evaluation metrics and techniques used in IMDb score prediction:

### **1. Mean Absolute Error (MAE):**

- MAE measures the average absolute difference between the predicted IMDb scores and the actual IMDb scores. A lower MAE indicates better performance.

### **2. Mean Squared Error (MSE):**

- MSE calculates the average of the squared differences between predicted and actual IMDb scores. It penalizes larger errors more heavily. Smaller MSE values are better.

### **3. Root Mean Squared Error (RMSE):**

- RMSE is the square root of MSE, and it provides a measure of the average prediction error in the same units as the IMDb scores. Lower RMSE values indicate better performance.

### **4. R-squared (R<sup>2</sup>) Score:**

- R<sup>2</sup> measures the proportion of the variance in IMDb scores that is explained by the model. It ranges from 0 to 1, with higher values indicating better model fit. An R<sup>2</sup> of 1 means the model perfectly predicts IMDb scores.

### **5.Coefficient of Determination (COD):**

- COD, similar to  $R^2$ , measures how well the model explains the variance in IMDb scores. It can help assess the goodness of fit.

### **6.Adjusted R-squared:**

- Adjusted R-squared takes into account the number of predictors in the model and adjusts  $R^2$  accordingly. It helps prevent overfitting by penalizing the inclusion of unnecessary features.

### **7.Percentile Ranking:**

- You can calculate the percentile rank of the model's predictions to determine how often it correctly predicts IMDb scores compared to other movies.

### **8.Residual Analysis:**

- Analyze the model's residuals (the differences between predicted and actual scores). Plotting residuals can help identify patterns or heteroscedasticity, which can indicate model deficiencies.

### **9. Cross-Validation:**

- Use cross-validation techniques (e.g., k-fold cross-validation) to assess how well the model generalizes to unseen data. Cross-validation provides a more robust estimate of the model's performance.

### **10.Bias-Variance Tradeoff:**

Consider the balance between bias and variance. High bias indicates underfitting, while high variance indicates overfitting. Achieving the right balance is crucial for model performance.

# IMDB SCORE PREDICTION IMPLEMENTATION

The task is to build a model for predicting the IMDB rating for the various movies

## Importing the required library for the task

In [1]:

```
import numpy as np
import pandas as pd
import csv
from scipy import stats
import matplotlib.pyplot as plt
```

In [2]:

```
rawdata = pd.read_csv('movie_metadata.csv')
```

In [3]:

```
rawdata.head()
```

In [4]:

```
rawdata.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5043 entries, 0 to 5042
```

```
Data columns (total 29 columns):
```

#	Column	Non-Null Count	Dtype
0	color	5024 non-null	object
1	director_name	4939 non-null	object
2	num_critic_for_reviews	4993 non-null	float64
3	duration	5028 non-null	float64
4	director_facebook_likes	4939 non-null	float64
5	actor_3_facebook_likes	5020 non-null	float64
6	actor_2_name	5030 non-null	object
7	actor_1_facebook_likes	5036 non-null	float64
8	gross	4159 non-null	float64
9	genres	5043 non-null	object
10	actor_1_name	5036 non-null	object
11	movie_title	5043 non-null	object
12	num_voted_users	5043 non-null	int64
13	cast_total_facebook_likes	5043 non-null	int64
14	actor_3_name	5020 non-null	object
15	facenumber_in_poster	5030 non-null	float64
16	plot_keywords	4890 non-null	object
17	movie_imdb_link	5043 non-null	object
18	num_user_for_reviews	5022 non-null	float64
19	language	5029 non-null	object
20	country	5038 non-null	object
21	content_rating	4740 non-null	object
22	budget	4551 non-null	float64
23	title_year	4935 non-null	float64
24	actor_2_facebook_likes	5030 non-null	float64
25	imdb_score	5043 non-null	float64
26	aspect_ratio	4714 non-null	float64
27	movie_facebook_likes	5043 non-null	int64

```
28 Unnamed: 28          0 non-null   float64
dtypes: float64(14), int64(3), object(12)
memory usage: 1.1+ MB
```

**There are many variables in this dataset which account of integers, float values and categorical values. All these variables are not useful while considering to build a model for determining the IMDB ratings for the movies. These variables need to be analyzed for their usefulness.**

In [5]:

```
rawdata.isnull().sum(axis = 0)
```

Out[5]:

```
color                19
director_name        104
num_critic_for_reviews    50
duration             15
director_facebook_likes  104
actor_3_facebook_likes   23
actor_2_name         13
actor_1_facebook_likes    7
gross               884
genres               0
actor_1_name          7
movie_title           0
num_voted_users        0
cast_total_facebook_likes  0
actor_3_name          23
facenumber_in_poster    13
plot_keywords         153
movie_imdb_link         0
num_user_for_reviews    21
language             14
country              5
content_rating        303
budget              492
title_year           108
actor_2_facebook_likes   13
imdb_score            0
aspect_ratio          329
movie_facebook_likes     0
Unnamed: 28           5043
dtype: int64
```

**Finding out the null values in the dataset help us to find anomalies which can be taken care of as they would not help to create good models for continuous predictions.**


**Histograms allow us to figure out if the data is normally distributed and not skewed for particular values. Below is the histogram for the IMDB scores. We can see that there are many movies which have scores between the range of 5 and 8, which resembles a good chunk of data.**

In [6]:

```
figg = rawdata['imdb_score'].hist(bins=80, figsize=[14,6])
figg.plot()
```

Out[6]:

[]



## Finding the standard deviation and variance of the scores for the movies

```
rawdata['imdb_score'].std()
```

In [7]:

```
1.125115865732819
```

Out[7]:

```
rawdata['imdb_score'].var()
```

In [8]:

```
1.2658857113237107
```

Out[8]:

## Analysis of the data to find various information

### Finding the genres and top 10 countries for movies with IMDB scores greater than 7

```
country = rawdata['country'].value_counts()
goodmovies = rawdata.loc[rawdata['imdb_score'] >= 7]
genrecountry = goodmovies.groupby('genres')['country'].count()
ascgenre = genrecountry.sort_values(ascending=False)
top10country = country[:10]
```

In [9]:

```
ascgenre
```

In [10]:

```
genres
Drama                132
Drama|Romance         82
Comedy|Drama          73
Comedy|Drama|Romance  59
Crime|Drama|Thriller  47
```


Out[10]:

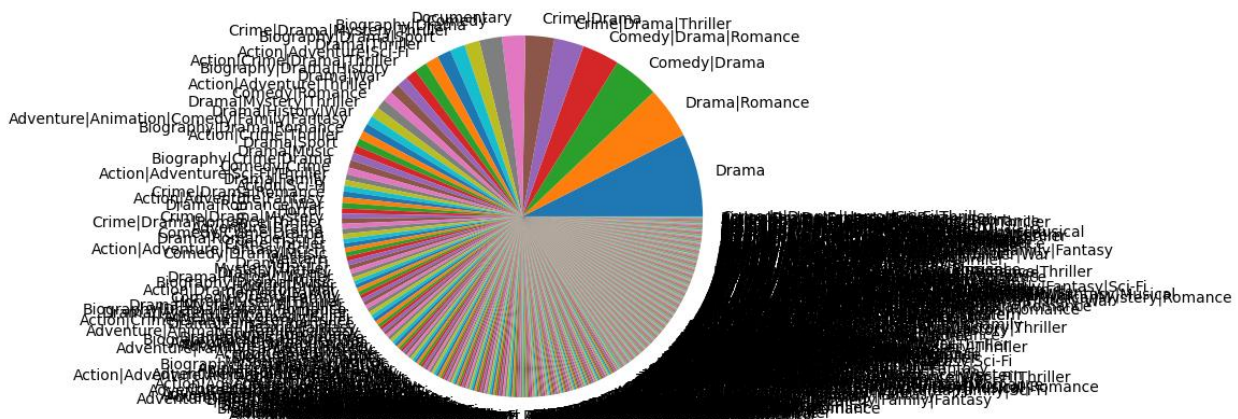
```
...
Action|Drama|Fantasy|Romance      1
Biography|Documentary|Sport       1
Action|Drama|Fantasy|Sci-Fi       1
Comedy|Drama|Horror|Sci-Fi|Thriller 1
Adventure|Drama|Fantasy|Mystery    0
Name: country, Length: 501, dtype: int64
```

In [11]:

```
top10country.plot(kind = 'bar')
plt.show()
fig1, ax1 = plt.subplots()
ascgenre.plot(kind = 'pie')
labels = 'Action', 'Adventure', 'Drama', 'Animation', 'Comedy', 'Mystery', 'Crime', 'Biography', 'Fantasy',
'Documentary', 'Sci-Fi', 'Horror', 'Romance', 'Family', 'Western', 'Musical'

ax1.axis('equal')
plt.show()
```





**From the above charts we can figure out that USA produced the most films based of the data available followed by other countries in the list. The pie chart explains the different genres of films that had IMDB scores of more than 7. Action, adventure and drama lead other genres and are considered more popular for good movie types.**

### Finding the counts of film ratings for each and every country in the dataset.

In [12]:

```
from pandas import DataFrame
countryrating = rawdata.groupby(['country','content_rating']).size()
rating = DataFrame(countryrating)
rating
```

Out[12]:

0

country	content_rating	
Afghanistan	PG-13	1
Argentina	R	3
	Unrated	1
Aruba	R	1
Australia	G	2
...	...	...



0

country	content_rating	
USA	Unrated	38
	X	12
	M	1
West Germany	PG	1
	R	1

176 rows  $\times$  1 columns

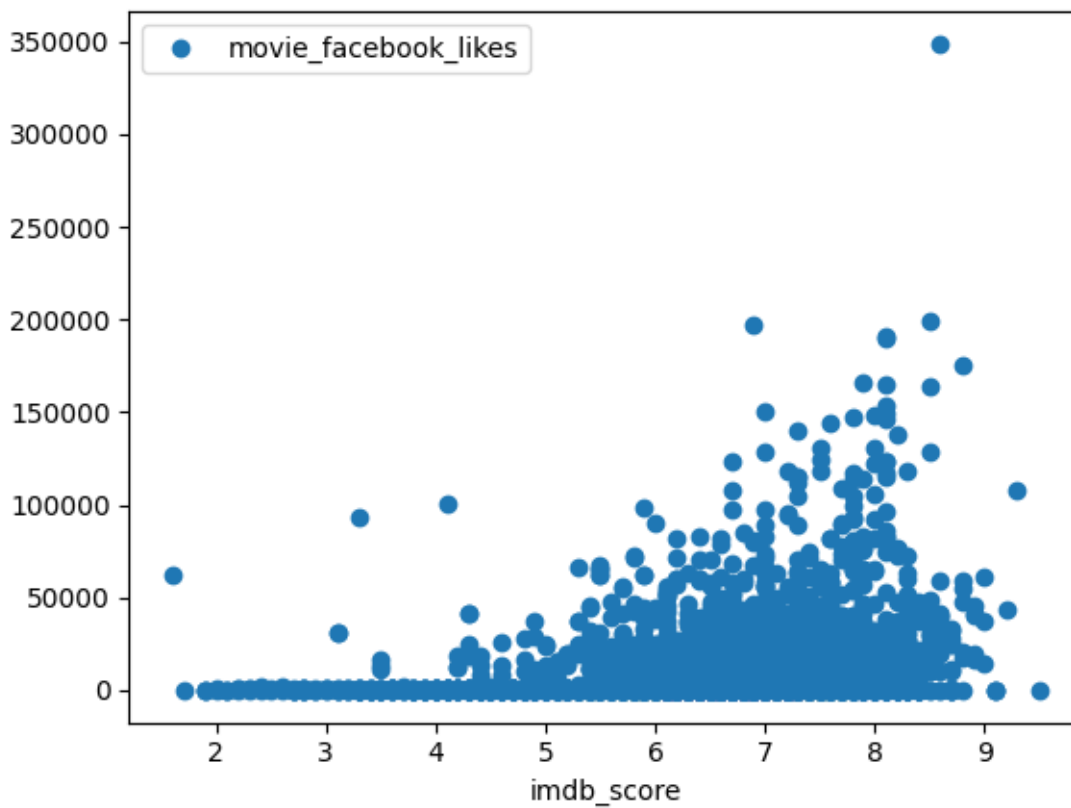
## Finding relationship between IMDB score and facebook likes

In [13]:

```
rawdata.plot(x='imdb_score', y='movie_facebook_likes', style='o')
```

Out[13]:

<Axes: xlabel='imdb\_score'>



## Processing Data for building machine learning models to predict IMDB ratings which are a series of continuous values and not classes

In [14]:

```
newdata =  
rawdata[['color','num_critic_for_reviews','duration','director_facebook_likes','actor_1_facebook_likes','actor_2_face  
book_likes','actor_3_facebook_likes','gross','genres','movie_title','num_voted_users','cast_total_facebook_likes','face  
number_in_poster','num_user_for_reviews','language','country','content_rating','budget','title_year','aspect_ratio','mo  
vie_facebook_likes','imdb_score']]
```

In [15]:

```
newdata[:5]
```

Out[15]:

## Null values have to be taken care of before feeding it to the model

In [16]:

```
newdata = newdata.fillna(value=0)
```

In [17]:

```
newdata.isnull().sum(axis = 0)
```

Out[17]:

```
color          0  
num_critic_for_reviews    0  
duration        0  
director_facebook_likes  0  
actor_1_facebook_likes   0  
actor_2_facebook_likes   0  
actor_3_facebook_likes   0  
gross           0  
genres          0  
movie_title      0  
num_voted_users    0  
cast_total_facebook_likes  0  
facenumber_in_poster  0  
num_user_for_reviews    0  
language         0  
country          0  
content_rating     0  
budget           0  
title_year        0  
aspect_ratio      0  
movie_facebook_likes    0  
imdb_score        0  
dtype: int64
```

## Writing a function to convert all the categorical variables in the dataset to numeric values

In [18]:

```
class MultiColumnLabelEncoder:  
    def __init__(self, columns = None):  
        self.columns = columns # array of column names to encode  
  
    def fit(self, X, y=None):
```

```
return self # not relevant here
```

```
def transform(self,X):  
    """  
    Transforms columns of X specified in self.columns using  
    LabelEncoder(). If no columns specified, transforms all  
    columns in X.  
    """  
    output = X.copy()  
    if self.columns is not None:  
        for col in self.columns:  
            output[col] = LabelEncoder().fit_transform(output[col])  
    else:  
        for colname,col in output.iteritems():  
            output[colname] = LabelEncoder().fit_transform(col)  
    return output  
  
def fit_transform(self,X,y=None):  
    return self.fit(X,y).transform(X)
```


In [19]:

```
from sklearn.preprocessing import LabelEncoder  
from sklearn.pipeline import Pipeline  
newdata = MultiColumnLabelEncoder(columns =  
['color','num_critic_for_reviews','duration','director_facebook_likes','actor_1_facebook_likes','actor_2_facebook_likes',  
'actor_3_facebook_likes','gross','genres','movie_title','num_voted_users','cast_total_facebook_likes','facenumber_in_poster',  
'num_user_for_reviews','language','country','content_rating','budget','title_year','aspect_ratio','movie_facebook_likes',  
'imdb_score']).fit_transform((newdata).astype(str))
```

**One of the main task herre is to find correlations between the variables in the dataset. If there are variables that are closely related to each other it would be a redundant variable which would be taking up resources during training and might not yield a good model for predicting the scores.**

In [20]:

```
from matplotlib import cm as cm  
import pandas  
import numpy  
  
names =  
['color','num_critic_for_reviews','duration','director_facebook_likes','actor_1_facebook_likes','actor_2_facebook_likes',  
'actor_3_facebook_likes','gross','genres','movie_title','num_voted_users','cast_total_facebook_likes','facenumber_in_poster',  
'num_user_for_reviews','language','country','content_rating','budget','title_year','aspect_ratio','movie_facebook_likes',  
'imdb_score']  
correlations = newdata.corr()  
# plot correlation matrix  
fig = plt.figure(figsize=(20,10))  
ax = fig.add_subplot(111)  
cax = ax.matshow(correlations,cmap=plt.cm.Blues)  
fig.colorbar(cax)  
ticks = numpy.arange(0,22,1)  
ax.set_xticks(ticks)  
ax.set_yticks(ticks)  
ax.set_xticklabels(names)  
ax.set_yticklabels(names)  
plt.show()
```



**From the above correlation matrix we can see that the selected variable are not closely related to each other and can be used for training the models.**

## Splitting the data into train and test sets

In [21]:

```
from sklearn.model_selection import train_test_split
newdata = newdata.drop(columns=['imdb_score'])
finaldata = pd.concat([newdata, rawdata['imdb_score']], axis=1)
x = finaldata.loc[:, finaldata.columns != 'imdb_score']
y = finaldata['imdb_score']
```

In [22]:

```
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2)
print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)
(4034, 21) (4034,)
(1009, 21) (1009,)
```

## Building different models to figure out which performs the best for regression

### Simple linear regression model

In [23]:

```
from sklearn import linear_model
lm = linear_model.LinearRegression()
```

```
model = lm.fit(X_train, y_train)
y_pred = model.predict(X_test)
```

In [24]:

```
from sklearn import metrics
print(np.sqrt(metrics.mean_squared_error(y_test,y_pred)))
1.0417184593862583
```

**From the above results we can see that the model doesn't perform well to predict the exact scores on the test set. This needs to be improved and there are many ways to approach such problems. One of the problems is dimensionality reduction, where we get only those features which seem important for predicting.**

**Random forest for extracting features which seem important. Random forest being an ensemble model helps to select features based on selecting from many subsets. This helps to eliminate features which would be biased and give equal importance to all.**

In [25]:

```
from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor(random_state=1, max_depth=100)
rf.fit(X_train, y_train)
```


Out[25]:

```
RandomForestRegressor(max_depth=100, random_state=1)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook. On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [26]:

```
features = x.columns
importances = rf.feature_importances_
indices = np.argsort(importances)[:22] # top 10 features
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='b', align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



**From the above methods for feature selection we can see that 'duration' has come up as most important feature while determining the features to predict score for the movies. But duration isn't really a feature we can rely on and other variables has to be taken into account.**

In [27]:

```
ranfor =
finaldata[['duration','title_year','movie_facebook_likes','num_critic_for_reviews','num_user_for_reviews','genres','cast_total_facebook_likes','num_voted_users']]
```

In [28]:

```
X_train, X_test, y_train, y_test = train_test_split(ranfor, y, test_size=0.2)
print (X_train.shape, y_train.shape)
print (X_test.shape, y_test.shape)

(4034, 8) (4034,)
(1009, 8) (1009,)
```

## Random forest model

In [29]:

```
from sklearn.ensemble import RandomForestRegressor
rf1 = RandomForestRegressor(random_state=9, max_depth=100)
rf1.fit(X_train, y_train)
```

Out[29]:

```
RandomForestRegressor(max_depth=100, random_state=9)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook. On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [30]:

```
y_pred = rf1.predict(X_test)
```

In [31]:

```
from sklearn.metrics import r2_score
coefficient_of_determination = r2_score(y_test,y_pred)
print(coefficient_of_determination*100)
31.967307067121308
```

In [32]:

```
ssr = np.sum((y_pred - y_test)**2)
```

In [33]:

```
print(ssr)
893.7796450000002
```

**From the above result we can see that the model has a low accuracy because it is a regression problem and it's very difficult to predict continuous variables**

## Using PCA for dimensionality reduction instead of Random Forest

In [34]:

```
from sklearn.decomposition import PCA
pca = PCA(n_components=4)
pca_result = pca.fit_transform(ranfor)
```

In [35]:

```
X_train, X_test, y_train, y_test = train_test_split(pca_result, y, test_size=0.2)
print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)
(4034, 4) (4034,)
(1009, 4) (1009,)
```

In [36]:

```
from sklearn import linear_model
lm = linear_model.LinearRegression()
model1 = lm.fit(X_train, y_train)
y_pred = model1.predict(X_test)
```

In [37]:

```
from sklearn.ensemble import RandomForestRegressor
rf2 = RandomForestRegressor(random_state=5, max_depth=1000)
rf2.fit(X_train, y_train)
```

Out[37]:

```
RandomForestRegressor(max_depth=1000, random_state=5)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook. On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [38]:

```
y_pred = rf2.predict(X_test)
```

## RMSE vs R2

In [39]:

```
from sklearn import metrics
print(np.sqrt(metrics.mean_squared_error(y_test,y_pred)))
1.1272981104093134
```

In [40]:

```
from sklearn.metrics import r2_score
coefficient_of_determination = r2_score(y_test,y_pred)
print(coefficient_of_determination*100)
```

1.37188201899473

## Building a XGBoost model for prediction

In [41]:

```
import xgboost as xgb
xg_reg = xgb.XGBRegressor(objective='reg:linear', colsample_bytree = 0.3, learning_rate = 0.1, max_depth = 5,
alpha = 10, n_estimators = 100)
xg_reg.fit(X_train,y_train)
preds = xg_reg.predict(X_test)

[23:01:28] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-07593ffd91cd9da33-1\xgboost\xgboost-ci-windows\src\objective\regression_obj.cu:213: reg:linear is now deprecated in favor of reg:squarederror.
```

In [42]:

```
coefficient_of_determination = r2_score(y_test,preds)
print(coefficient_of_determination*100)
print(np.sqrt(metrics.mean_squared_error(y_test,preds)))

1.686629884237012
1.125497922947681
```

**We can see that results from XGBoost has slightly increased and still isn't best for predicting the right scores and can be improved using cross validation**

## Applying Cross validation with XGBoost to build a more robust model than the previous ones

In [43]:

```
data_dmatrix = xgb.DMatrix(data=x,label=y)
```

In [44]:

```
params = {"objective": "reg:linear", 'colsample_bytree': 0.3, 'learning_rate': 0.1,
          'max_depth': 5, 'alpha': 10}

cv_results = xgb.cv(dtrain=data_dmatrix, params=params, nfold=4,
                    num_boost_round=50, early_stopping_rounds=10, metrics="rmse", as_pandas=True, seed=123)

[23:01:28] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-07593ffd91cd9da33-1\xgboost\xgboost-ci-windows\src\objective\regression_obj.cu:213: reg:linear is now deprecated in favor of reg:squarederror.
[23:01:28] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-07593ffd91cd9da33-1\xgboost\xgboost-ci-windows\src\objective\regression_obj.cu:213: reg:linear is now deprecated in favor of reg:squarederror.
[23:01:28] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-07593ffd91cd9da33-1\xgboost\xgboost-ci-windows\src\objective\regression_obj.cu:213: reg:linear is now deprecated in favor of reg:squarederror.
[23:01:28] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-07593ffd91cd9da33-1\xgboost\xgboost-ci-windows\src\objective\regression_obj.cu:213: reg:linear is now deprecated in favor of reg:squarederror.
```

In [45]:

```
cv_results.head()
```

Out[45]:



	train-rmse-mean	train-rmse-std	test-rmse-mean	test-rmse-std
0	5.463068	0.014895	5.463214	0.049900
1	4.940707	0.013058	4.940953	0.051221
2	4.472980	0.011846	4.473388	0.052249
3	4.053943	0.010850	4.053918	0.053153
4	3.678052	0.009804	3.677957	0.052937

In [46]:

```
print((cv_results["test-rmse-mean"]).tail(1))
```

```
49 0.944848
```

```
Name: test-rmse-mean, dtype: float64
```

In [47]:

```
xg_reg = xgb.train(params=params, dtrain=data_dmatrix, num_boost_round=10)
```

```
[23:01:29] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-07593ffd91cd9da33-1\xgboost\xgboost-ci-windows\src\objective\regression_obj.cu:213: reg:linear is now deprecated in favor of reg:squarederror.
```

## Conclusion

We can see that the RMSE value for predicting has improved significantly from the previous results thanks to Cross validation. This can further be improved by fine tuning the parameters and improving the model performance. For now the XGBoost model with cross validation helps to predict the rating best.

The other way to improve the predictions is to convert the imdb\_scores from continious values to classes. This way the problem gets converted into a classification problem with 10 different classes and classification algorithms can be used to predict the imdb score classes.